

```
import os
import torch
import cv2
import numpy as np
import math
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from tensorflow.keras.datasets import cifar10
```

```
2024-10-04 18:46:40.415250: I tensorflow/core/util/port.cc:111] oneDNN custom operations are on. You may see slightly different numer
2024-10-04 18:46:40.449528: I tensorflow/tsl/cuda/cudart_stub.cc:28] Could not find cuda drivers on your machine, GPU will not be use
2024-10-04 18:46:40.626665: E tensorflow/compiler/xla/stream_executor/cuda/cuda_dnn.cc:9342] Unable to register cuDNN factory: Attempt
2024-10-04 18:46:40.626748: E tensorflow/compiler/xla/stream_executor/cuda/cuda_fft.cc:609] Unable to register cuFFT factory: Attempt
2024-10-04 18:46:40.627627: E tensorflow/compiler/xla/stream_executor/cuda/cuda_blas.cc:1518] Unable to register cuBLAS factory: Atte
2024-10-04 18:46:40.694574: I tensorflow/core/platform/cpu_feature_guard.cc:182] This TensorFlow binary is optimized to use available
To enable the following instructions: AVX2 AVX_VNNI FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
2024-10-04 18:46:42.587961: W tensorflow/compiler/tf2tensorrt/utils/py_utils.cc:38] TF-TRT Warning: Could not find TensorRT
```

Exercise 1

```
np_arr = np.random.randint(3, 10, size=(3, 4))
print(np_arr, type(np_arr))
```

```
[[8 5 3 7]
 [4 3 6 3]
 [7 9 8 5]] <class 'numpy.ndarray'>
```

```
print(np_arr[0][3])
```

```
7
```

```
print(np_arr[:, :2])
```

```
[[8 5]
 [4 3]
 [7 9]]
```

```
np_arr_1 = np.random.randint(3, 10, size=12)
print(np_arr_1)
```

```
[4 3 5 6 3 8 4 3 7 9 9 9]
```

```
np_arr_1_reshaped = np_arr_1.reshape((3,4))
print(np_arr_1_reshaped)
```

```
[[4 3 5 6]
 [3 8 4 3]
 [7 9 9 9]]
```

```
arr1 = np.random.randint(1,10, size=(3,3))
arr2 = np.random.randint(1,10, size=(3,3))
```

```
print(arr1@arr2)
```

```
[[ 74  71  75]
 [ 97 128  90]
 [ 44  55  51]]
```

```
print("Mean: ",arr1.mean())
print("Standard deviation: ",arr1.std())
print("Median: ",np.median(arr1))
```

```
Mean: 4.0
Standard deviation: 2.3570226039551585
Median: 4.0
```

```
print("Before: ", np_arr, "\n")
print("After: ",np_arr.flatten())
```

```

➦ Before:  [[8 5 3 7]
             [4 3 6 3]
             [7 9 8 5]]

After:  [8 5 3 7 4 3 6 3 7 9 8 5]

```

```
array1 = np.array([[1, 2],
                  [3, 4]])
```

```
array2 = np.array([[5, 6],
                  [7, 8]])
```

```
vertical_stack = np.vstack((array1, array2))
print("Vertical Stacking:\n", vertical_stack)
```

```
horizontal_stack = np.hstack((array1, array2))
print("Horizontal Stacking:\n", horizontal_stack)
```

```

➦ Vertical Stacking:
[[1 2]
 [3 4]
 [5 6]
 [7 8]]
Horizontal Stacking:
[[1 2 5 6]
 [3 4 7 8]]

```

```

n=10
random_array = np.random.normal(loc=0.0, scale=1.0, size=n)
other_array = np.random.normal(loc=1.0, scale=1.0, size=n)
addition = random_array + other_array
subtraction = random_array - other_array
multiplication = random_array * other_array
division = random_array / other_array

```

```

print("Random Array:\n", random_array)
print("Other Array:\n", other_array)
print("Element-wise Addition:\n", addition)
print("Element-wise Subtraction:\n", subtraction)
print("Element-wise Multiplication:\n", multiplication)
print("Element-wise Division:\n", division)

```

```

➦ Random Array:
[-0.06877117  0.01451558 -0.24019468 -0.69013421  0.32595183  0.80754994
 -1.38955979 -0.5016213  -0.68610501 -2.20082033]
Other Array:
[ 0.37875933  3.18767661  1.53419819 -0.12697382  2.56237615  0.94156057
  1.64885298  1.36727634 -0.51365855  1.25106766]
Element-wise Addition:
[ 0.30998816  3.2021922  1.29400352 -0.81710803  2.88832798  1.74911051
  0.25929319  0.86565504 -1.19976356 -0.94975267]
Element-wise Subtraction:
[-0.4475305  -3.17316103 -1.77439287 -0.56316039 -2.23642431 -0.13401064
 -3.03841277 -1.86889764 -0.17244646 -3.45188799]
Element-wise Multiplication:
[-0.02604772  0.04627099 -0.36850624  0.08762898  0.8352112  0.76035718
 -2.29117979 -0.68585494  0.35242371 -2.75337514]
Element-wise Division:
[-1.81569570e-01  4.55365621e-03 -1.56560396e-01  5.43524795e+00
  1.27206864e-01  8.57671782e-01 -8.42743293e-01 -3.66876310e-01
  1.33572197e+00 -1.75915372e+00]

```


✓ Exercise 2

```

img = cv2.imread("doggo.jpeg")
plt.imshow(img)

```

 <matplotlib.image.AxesImage at 0x7fcee1ff4c0>



```
image = cv2.imread("doggo.jpeg")
image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
height, width = image.shape[:2]

resized_image = cv2.resize(image, (width // 2, height // 2))
print(resized_image.shape)
```

 (96, 102, 3)

```
grey_image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
plt.imshow(grey_image, cmap="gray")
```

 <matplotlib.image.AxesImage at 0x7fcee0fea10>

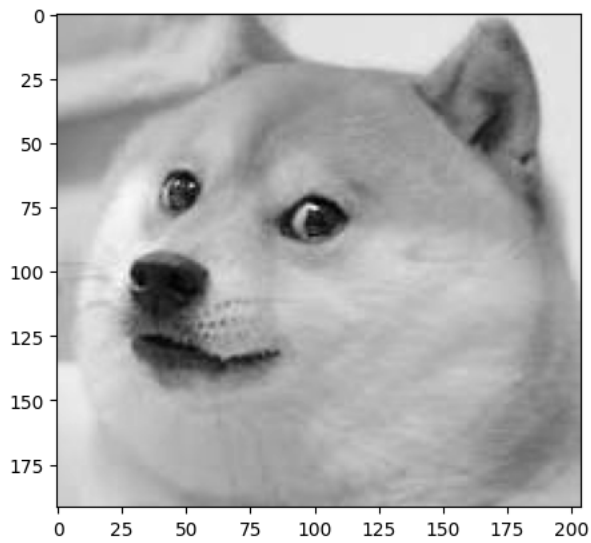


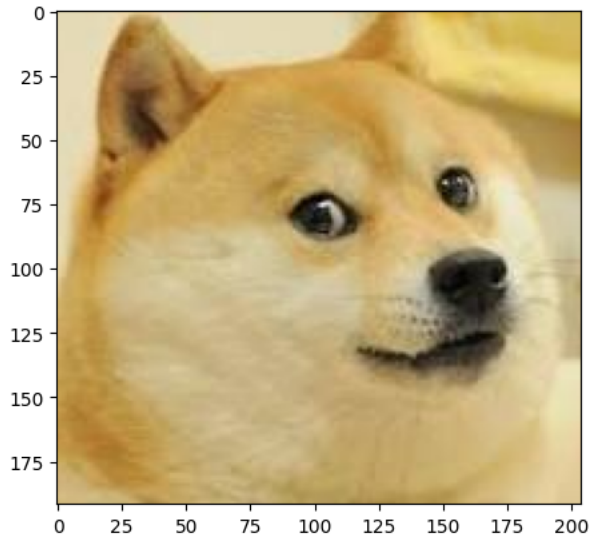
image.shape

 (192, 204, 3)

```
horizontal_flip = cv2.flip(image, 1)
vertical_flip = cv2.flip(image, 0)

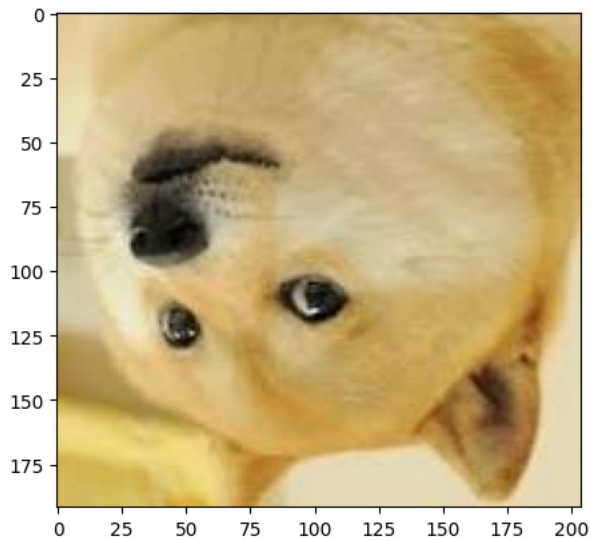
plt.imshow(horizontal_flip)
```

 <matplotlib.image.AxesImage at 0x7fcee90cada0>




```
plt.imshow(vertical_flip)
```

 <matplotlib.image.AxesImage at 0x7fcee8f3f070>



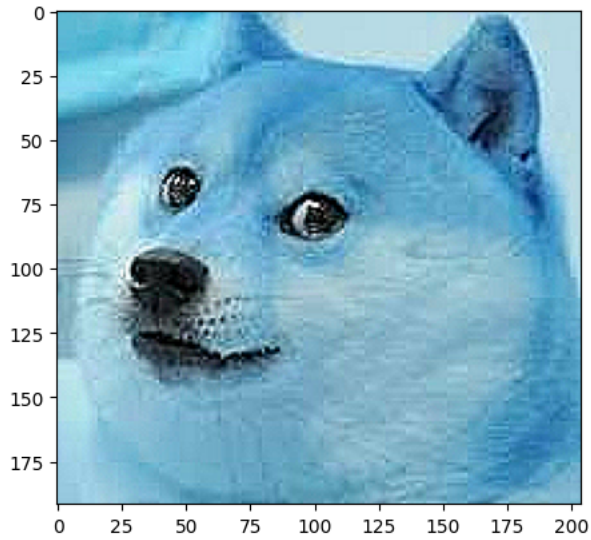
```
smoothing_kernel = np.ones((5, 5), np.float32) / 25
smoothed_img = cv2.filter2D(image, -1, smoothing_kernel)
smoothed_img_rgb = cv2.cvtColor(smoothed_img, cv2.COLOR_BGR2RGB)
plt.imshow(smoothed_img_rgb)
```

 <matplotlib.image.AxesImage at 0x7fcee8403bb0>



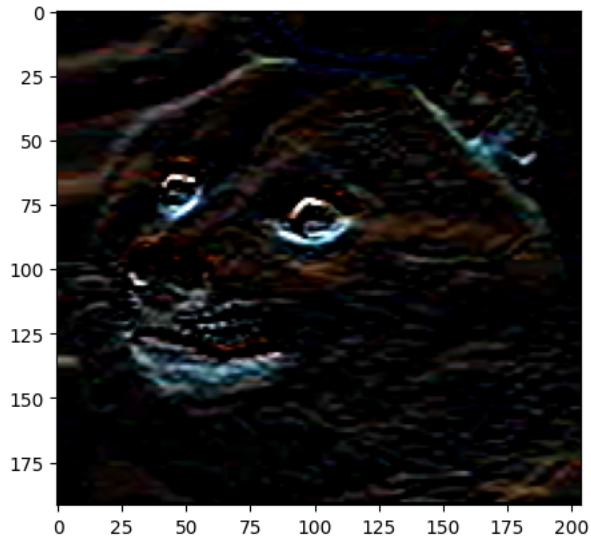
```
sharpening_kernel = np.array([[0, -1, 0], [-1, 5, -1], [0, -1, 0]])
sharp_img = cv2.filter2D(image, -1, sharpening_kernel)
sharp_img_rgb = cv2.cvtColor(sharp_img, cv2.COLOR_BGR2RGB)
plt.imshow(sharp_img_rgb)
```

 <matplotlib.image.AxesImage at 0x7fcee8326260>



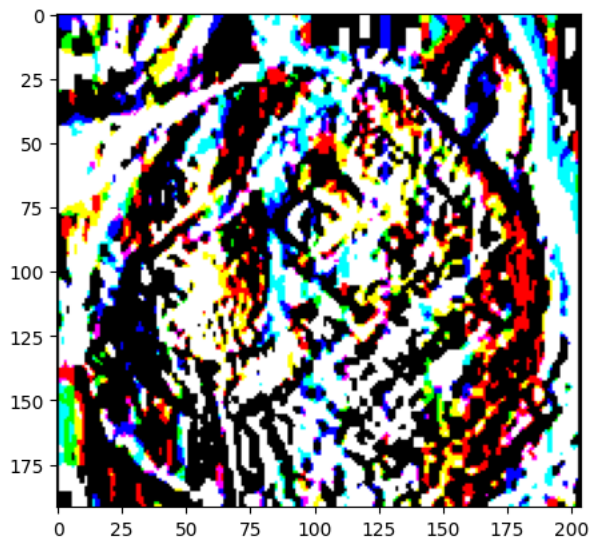
```
SVkernel = np.array([[ -1, -2, -1], [ 0,  0,  0], [ 1,  2,  1]])
edge_img = cv2.filter2D(image, -1, SVkernel)
edge_img_rgb = cv2.cvtColor(edge_img, cv2.COLOR_BGR2RGB)
plt.imshow(edge_img_rgb)
```

 <matplotlib.image.AxesImage at 0x7fcee847a200>



```
sobelx = cv2.Sobel(src=image, ddepth=cv2.CV_64F, dx=1, dy=0, ksize=5)
plt.imshow(sobelx)
```

 Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
<matplotlib.image.AxesImage at 0x7fcee90ae500>



```
A = torch.rand(3, 3, requires_grad=True)
B = torch.rand(3, 3)
```

```
C = torch.matmul(A, B)
```

```
C.backward(torch.ones_like(C))
```

```
print("Gradient of A:\n", A.grad)
```

```
↪ Gradient of A:
  tensor([[0.8512, 0.8763, 1.4656],
          [0.8512, 0.8763, 1.4656],
          [0.8512, 0.8763, 1.4656]])
```

```
tensor1 = torch.rand(3, 1)
tensor2 = torch.rand(1, 3)
```

```
broadcast_result = tensor1 + tensor2
C = torch.rand(3, 3)
result = broadcast_result * C
```

```
print("Result of broadcasting:\n", result)
```

```
↪ Result of broadcasting:
  tensor([[1.1773, 0.9084, 0.6238],
          [0.2204, 0.4859, 0.2680],
          [0.6087, 0.2442, 0.9910]])
```

```
tensor = torch.rand(6, 4)
```

```
reshaped_tensor = tensor.view(3, 8)
```

```
sliced_tensor = reshaped_tensor[:, :2]
```

```
print("Reshaped Tensor:\n", reshaped_tensor)
print("Sliced Tensor (first two columns):\n", sliced_tensor)
```

```
↪ Reshaped Tensor:
  tensor([[0.0644, 0.0934, 0.6452, 0.5972, 0.9199, 0.5601, 0.4757, 0.7783],
          [0.8037, 0.8217, 0.3071, 0.2832, 0.0558, 0.9997, 0.0272, 0.7795],
          [0.4950, 0.2929, 0.9127, 0.0625, 0.4126, 0.2511, 0.0801, 0.7203]])
Sliced Tensor (first two columns):
  tensor([[0.0644, 0.0934],
          [0.8037, 0.8217],
          [0.4950, 0.2929]])
```

```
np_array = np.random.rand(3, 3)
```

```
torch_tensor = torch.from_numpy(np_array)
```

```
modified_tensor = torch_tensor * 3
```

```
modified_np_array = modified_tensor.numpy()
```

```
print("Original NumPy array:\n", np_array)
print("Modified NumPy array:\n", modified_np_array)
```

```
↪ Original NumPy array:
  [[0.32267746 0.71822701 0.21827624]
  [0.45957789 0.24949688 0.52984562]
  [0.43003467 0.68085862 0.71411156 ]]
Modified NumPy array:
  [[0.96803238 2.15468104 0.65482871]
  [1.37873367 0.74849064 1.58953685]
  [1.29010402 2.04257586 2.14234679]]
```

```
tensor_uniform = torch.rand(5, 5)
```

```
tensor_normal = torch.randn(5, 5)
```

```
elementwise_product = tensor_uniform * tensor_normal
```

```
mean_val = elementwise_product.mean()
std_val = elementwise_product.std()
```

```
reshaped_tensor = elementwise_product.view(25)
sum_val = reshaped_tensor.sum()
```

```
print("Mean of result tensor:", mean_val.item())
print("Standard deviation of result tensor:", std_val.item())
print("Sum of all elements:", sum_val.item())
```

```
↪ Mean of result tensor: 0.027424873784184456
  Standard deviation of result tensor: 0.5586299300193787
```

Sum of all elements: 0.6856218576431274

Exercise 4

```
def sigmoid_math(x):
    return 1 / (1 + math.exp(-x))

def sigmoid_np(x):
    return 1 / (1 + np.exp(-x))

def sigmoid_grad(x):
    s = sigmoid_np(x)
    return s * (1 - s)

def image2vector(image):
    return image.reshape(-1, 1)

def normalizeRows(x):
    return x / np.linalg.norm(x, axis=1, keepdims=True)

def L1_loss(y_hat, y):
    return np.sum(np.abs(y_hat - y))

def L2_loss(y_hat, y):
    return np.sum((y_hat - y) ** 2)
```

```
x = 1.5
image = np.random.rand(3, 3, 3)
y_hat = np.array([0.9, 0.2, 0.4])
y = np.array([1.0, 0.0, 0.5])
matrix_x = np.array([[4, 3], [1, 2]])

print("Sigmoid using math.exp:", sigmoid_math(x))
print("Sigmoid using np.exp:", sigmoid_np(x))
print("Sigmoid gradient:", sigmoid_grad(x))
print("Image to vector:", image2vector(image))
print("Normalized rows:\n", normalizeRows(matrix_x))
print("L1 Loss:", L1_loss(y_hat, y))
print("L2 Loss:", L2_loss(y_hat, y))
```

↗ Sigmoid using math.exp: 0.8175744761936437

Sigmoid using np.exp: 0.8175744761936437

Sigmoid gradient: 0.14914645207033286

Image to vector: [[0.86429375]

[0.56291403]

[0.57270959]

[0.04184646]

[0.06848765]

[0.11637612]

[0.69159536]

[0.82905832]

[0.51452064]

[0.20064762]

[0.13402107]

[0.57797104]

[0.67334416]

[0.19806764]

[0.0221989]

[0.03247737]

[0.83618819]

[0.99573084]

[0.94705162]

[0.84372425]

[0.9256295]

[0.49159989]

[0.29451266]

[0.92984915]

[0.6152323]

[0.92049842]

[0.58456849]]

Normalized rows:

[[0.8 0.6

[0.4472136 0.89442719]]

L1 Loss: 0.39999999999999997

L2 Loss: 0.06

To do this locally, I will be using the CIFAR10 dataset from the keras datasets. From the 10 classes, I will be treating the cats class(label: 3) as the positive class, and the other 9 classes as negative classes.

```
(X_train, y_train), (X_test, y_test) = cifar10.load_data()
```

```
↳ Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz  
170498071/170498071 [=====] - 199s 1us/step
```

```
y_train_binary = np.where(y_train == 3, 1, 0).reshape(-1)  
y_test_binary = np.where(y_test == 3, 1, 0).reshape(-1)
```


```
X_train_flatten = X_train.reshape(X_train.shape[0], -1)  
X_test_flatten = X_test.reshape(X_test.shape[0], -1)
```

```
first_image = X_train[0]  
print(f"Shape of the first image (should be 32x32x3 for RGB): {first_image.shape}")
```

```
plt.imshow(first_image)  
plt.title("First Image (Example from CIFAR-10)")  
plt.show()
```

```
↳ Shape of the first image (should be 32x32x3 for RGB): (32, 32, 3)
```

First Image (Example from CIFAR-10)



```
logistic_regression_model = LogisticRegression(max_iter=1000)  
logistic_regression_model.fit(X_train_flatten, y_train_binary)
```

```
y_pred = logistic_regression_model.predict(X_test_flatten)
```

```
↳ /home/the_architect/.local/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed to converge  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression