

✓ 22AIE304 Deep Learning Labsheet 5

Name: Aniketh Vijesh

Roll No: AM.EN.U4AIE22009



```
import torch
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

import torchvision.datasets as datasets
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim

# Setting up the datasets and Dataloaders

train_set = datasets.MNIST(
    root="./data", train=True, download=True, transform=transforms.ToTensor()
)
test_set = datasets.MNIST(
    root="./data", train=False, download=True, transform=transforms.ToTensor()
)

train_loader = DataLoader(dataset=train_set, batch_size=64, shuffle=True, num_workers=2)
test_loader = DataLoader(dataset=test_set, batch_size=64, shuffle=False, num_workers=2)

class FeedforwardNeuralNet(nn.Module):
    def __init__(self, input_size=784, hidden_size=128, output_size=10, num_hidden_layers=2, activation=nn.ReLU()):
        super(FeedforwardNeuralNet, self).__init__()

        self.fc1 = nn.Linear(input_size, hidden_size)

        self.hidden_layers = nn.ModuleList([nn.Linear(hidden_size, hidden_size) for _ in range(num_hidden_layers - 1)])

        self.output = nn.Linear(hidden_size, output_size)

        self.activation = activation

    def forward(self, x):
        x = x.view(x.size(0), -1)

        x = self.activation(self.fc1(x))

        for layer in self.hidden_layers:
            x = self.activation(layer(x))

        x = self.output(x)
        return x

# Training function
def train_model(model, train_loader, optimizer, criterion, num_epochs):
    model.train()
    train_losses = []
    train_accuracies = []
    for epoch in range(num_epochs):
        correct = 0
        total = 0
        epoch_loss = 0
        for images, labels in train_loader:
            optimizer.zero_grad()
            outputs = model(images)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            epoch_loss += loss.item()

        epoch_loss /= len(train_loader)
        train_losses.append(epoch_loss)

        for images, labels in train_loader:
            outputs = model(images)
            _, predicted = torch.max(outputs.data, 1)
            total += 1
            correct += predicted == labels
```

```

_, predicted = torch.max(outputs, 1)
total += labels.size(0)
correct += (predicted == labels).sum().item()

epoch_accuracy = correct / total
train_losses.append(epoch_loss / len(train_loader))
train_accuracies.append(epoch_accuracy)
print(f"Epoch {epoch+1}/{num_epochs}, Loss: {train_losses[-1]:.4f}, Accuracy: {epoch_accuracy:.4f}")

return train_losses, train_accuracies

```



```

# Plot training loss and accuracy
def plot_training_curves(train_losses, train_accuracies):
    epochs = range(1, len(train_losses) + 1)
    plt.figure(figsize=(12, 5))

    plt.subplot(1, 2, 1)
    plt.plot(epochs, train_losses, label="Training Loss")
    plt.xlabel("Epochs")
    plt.ylabel("Loss")
    plt.title("Training Loss Curve")
    plt.legend()

    plt.subplot(1, 2, 2)
    plt.plot(epochs, train_accuracies, label="Training Accuracy", color='orange')
    plt.xlabel("Epochs")
    plt.ylabel("Accuracy")
    plt.title("Training Accuracy Curve")
    plt.legend()

    plt.show()

# Evaluate the model on the test set
def evaluate_model(model, test_loader):
    model.eval()
    correct = 0
    total = 0
    all_labels = []
    all_predictions = []

    with torch.no_grad():
        for images, labels in test_loader:
            outputs = model(images)
            _, predicted = torch.max(outputs, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()
            all_labels.extend(labels.cpu().numpy())
            all_predictions.extend(predicted.cpu().numpy())

    accuracy = correct / total
    print(f"Test Accuracy: {accuracy:.4f}")

    # Confusion matrix
    cm = confusion_matrix(all_labels, all_predictions)
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=[i for i in range(10)])
    disp.plot(cmap=plt.cm.Blues)
    plt.title("Confusion Matrix")
    plt.show()

    return accuracy

```



```

activation_functions = {"ReLU": nn.ReLU(), "Sigmoid": nn.Sigmoid(), "Tanh": nn.Tanh(), "LeakyReLU": nn.LeakyReLU()}
results = {}

for name, activation_function in activation_functions.items():
    print(f"Training with {name} activation function...")
    model = FeedforwardNeuralNet(activation=activation_function)
    optimizer = optim.Adam(model.parameters(), lr=0.001)
    criterion = nn.CrossEntropyLoss()
    train_losses, train_accuracies = train_model(model, train_loader, optimizer, criterion, num_epochs=10)
    test_accuracy = evaluate_model(model, test_loader)
    results[name] = {"train_losses": train_losses, "train_accuracies": train_accuracies, "test_accuracy": test_accuracy}

```




```
plot_training_curves(train_losses, train_accuracies)
```



Training with ReLU activation function...

Epoch	Loss	Accuracy
1/10	0.3371	0.9079
2/10	0.1553	0.9548
3/10	0.1077	0.9685
4/10	0.0820	0.9760
5/10	0.0655	0.9802
6/10	0.0528	0.9842
7/10	0.0427	0.9872
8/10	0.0360	0.9893
9/10	0.0294	0.9916
10/10	0.0247	0.9932


Test Accuracy: 0.9777




Training with Sigmoid activation function...


Epoch	Loss	Accuracy
1/10	0.5352	0.8706
2/10	0.2272	0.9351
3/10	0.1745	0.9495
4/10	0.1419	0.9585
5/10	0.1179	0.9658
6/10	0.1001	0.9710
7/10	0.0855	0.9761
8/10	0.0735	0.9794
9/10	0.0642	0.9825
10/10	0.0563	0.9844

Test Accuracy: 0.9753






Training Loss Curve



Training Accuracy Curve




Training with Tanh activation function...


Epoch	Loss	Accuracy
1/10	0.3439	0.9050
2/10	0.1648	0.9525
3/10	0.1158	0.9661
4/10	0.0886	0.9745
5/10	0.0685	0.9809
6/10	0.0545	0.9845
7/10	0.0437	0.9881
8/10	0.0351	0.9910
9/10	0.0278	0.9932
10/10	0.0224	0.9952

Test Accuracy: 0.9772

Confusion Matrix



Training Loss Curve



Training Accuracy Curve

